

人工智慧專題

劉瓊如

October 7, 2022

1 Recommendation by graph

- PinSage
- Contrastive Learning

2 GCN

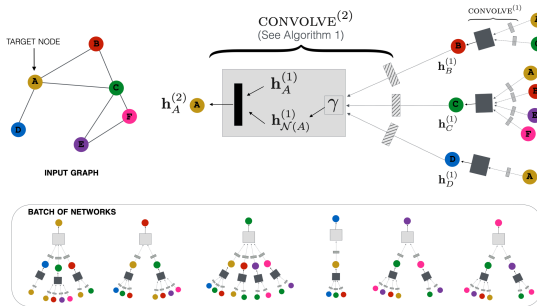
- Graph Theory
- GCMC

論文：Graph Convolutional Neural Networks for Web-Scale Recommender Systems

在 7.5 億的樣本裡，建立有 3 億個 nodes，18 億個邊的圖。PinSage 在 user studies, A/B test 裡，與單純的圖片相似度、語言相似度、混合模型、Pixie 都獲取較高的推薦。

算法解析：利用簡單的 GCN 為每個 pin 與 board 做 embedding，此 embedding 融合圖像與文字特徵來做高品質的 embedding。好的 embedding 不只在推薦或廣告投放有應用，用在 Pinterest 的實際服務上: related pins、搜尋、購物對下游任務非常重要，分類、聚類或重新排名。

PinSage 整體模型



整體模型：使用兩個 Convolve

PinSage Algorithm

Algorithm 1: CONVOLVE

Input : Current embedding \mathbf{z}_u for node u ; set of neighbor embeddings $\{\mathbf{z}_v | v \in \mathcal{N}(u)\}$, set of neighbor weights α ; symmetric vector function $\gamma(\cdot)$

Output: New embedding $\mathbf{z}_u^{\text{NEW}}$ for node u

- 1 $\mathbf{n}_u \leftarrow \gamma(\{\text{ReLU}(\mathbf{Q}\mathbf{h}_v + \mathbf{q}) \mid v \in \mathcal{N}(u)\}, \alpha)$;
 - 2 $\mathbf{z}_u^{\text{NEW}} \leftarrow \text{ReLU}(\mathbf{W} \cdot \text{CONCAT}(\mathbf{z}_u, \mathbf{n}_u) + \mathbf{w})$;
 - 3 $\mathbf{z}_u^{\text{NEW}} \leftarrow \mathbf{z}_u^{\text{NEW}} / \|\mathbf{z}_u^{\text{NEW}}\|_2$
-

$V = P \cup B$: 所有 nodes

- 將鄰居的鄰居向量經 affine transformation \rightarrow ReLU \rightarrow aggregation (importance pooling) \rightarrow 鄰居向量
- concatenate 自己與鄰居的向量 \rightarrow affine transformation \rightarrow ReLU \rightarrow 自己的向量
- normalized 成自己新的 embedding

PinSage Algorithm

- **Importance-based neighborhood**

For $u \in P$, $\mathcal{N}(u)$: 使用 random walk starting from u , 累計拜訪次數/random walk 長度 = 拜訪機率, 選取 T 個分數最高的機率的 pin 當成 u 的 neighborhood . 這樣會產生固定數量的鄰居。

Algorithm 1 的 γ : weighted mean (weight: 使用拜訪機率), 稱為 importance pooling . Pinterest 的 importance-based neighborhood 使用兩層結構 (two-fold), 也就是鄰居的鄰居來做自身的特徵抽取。

- **Stacking convolution**

Algorithm 1 是一層的 convolution

Algorithm 2 是疊加 convolution 使得可以得到 mini-batch 裡的每個 node 都可以有 k -convolution iteration 的 embedding , 層內的 Q, q, W, w 是分享權重, 但不同層的權重不同

PinSage Algorithm2

Algorithm 2: MINIBATCH

Input : Set of nodes $\mathcal{M} \subset \mathcal{V}$; depth parameter K ;
neighborhood function $\mathcal{N} : \mathcal{V} \rightarrow 2^{\mathcal{V}}$

Output: Embeddings $\mathbf{z}_u, \forall u \in \mathcal{M}$

/* Sampling neighborhoods of minibatch nodes. */

1 $\mathcal{S}^{(K)} \leftarrow \mathcal{M}$;

2 **for** $k = K, \dots, 1$ **do**

3 $\mathcal{S}^{(k-1)} \leftarrow \mathcal{S}^{(k)}$;

4 **for** $u \in \mathcal{S}^{(k)}$ **do**

5 $\mathcal{S}^{(k-1)} \leftarrow \mathcal{S}^{(k-1)} \cup \mathcal{N}(u)$;

6 **end**

7 **end**

/* Generating embeddings */

8 $\mathbf{h}_u^{(0)} \leftarrow \mathbf{x}_u, \forall u \in \mathcal{S}^{(0)}$;

9 **for** $k = 1, \dots, K$ **do**

10 **for** $u \in \mathcal{S}^{(k)}$ **do**

11 $\mathcal{H} \leftarrow \{\mathbf{h}_v^{(k-1)}, \forall v \in \mathcal{N}(u)\}$;

12 $\mathbf{h}_u^{(k)} \leftarrow \text{CONVOLVE}^{(k)}(\mathbf{h}_u^{(k-1)}, \mathcal{H})$

13 **end**

14 **end**

15 **for** $u \in \mathcal{M}$ **do**

16 $\mathbf{z}_u \leftarrow \mathbf{G}_2 \cdot \text{ReLU}(\mathbf{G}_1 \mathbf{h}_u^{(K)} + \mathbf{g})$

17 **end**

Loss function : Max-Margin ranking loss

利用 triplet loss 的概念，與正樣本拉近，負樣本拉遠到一個範圍外，但 triplet loss 使用的是歐式距離，這裡是內積。不過，作用的向量都是 unit vector，內積等於 cos 值。對於一個 query item z_q 與一個 candidate item z_i


$$J_G(z_q, z_i) = \mathbb{E}_{n_k \sim P_n(q)} \max\{0, z_q \cdot z_{n_k} - z_q \cdot z_i + \Delta\},$$

where $P_n(q)$ represents the distribution of negative examples for item q , Δ denotes the marginal hyper-parameter.

Note that $z_q \cdot z_{n_k} - z_q \cdot z_i + \Delta > 0$ 就會產生 loss，希望 $z_q \cdot z_{n_k} + \Delta \leq z_q \cdot z_i$ ，負樣本在一定的角度之外。

negative sampling：選取 500 個共用在每一個 minibatch

Angle margin



- If $\alpha = \cos^{-1}\left(\frac{z_i^T z_q}{|z_i||z_q|}\right) > \beta = \cos^{-1}\left(\frac{z_{n_k}^T z_q}{|z_{n_k}||z_q|}\right)$ or $\frac{z_i^T z_q}{|z_i||z_q|} < \frac{z_{n_k}^T z_q}{|z_{n_k}||z_q|}$, then produce loss
- 近似: $z_i^T z_q < z_{n_k}^T z_q$ 產生 loss
- $z_{n_k}^T z_q - z_i^T z_q > 0$ 產生 loss
- 在 margin Δ 範圍外 $z_{n_k}^T z_q + \Delta - z_i^T z_q > 0$ 產生 loss

Multi-GPU training with large minibatch

將一個 mini-batch 切分成等分 portion，每個 GPU 跑一個 portion，做 back propagation 時，同一個參數會在不同 portion 裡，累積 gradient 做一次更新。

假設 A, B node 在 portion 1 (batch 1), C, D node 在 portion 2 (batch 2), E, F node 在 portion 3 (batch 3)。

$$\hat{w}_{t+1} = w_t - \hat{\eta} \frac{1}{kn} \sum_{j < k} \sum_{x \in B_j} \nabla \ell(x, w_t)$$

在第 t 個 iteration 的參數 w_t ，共有 k 個 minibatch。

Curriculum training

抽樣 500 個 negative sample 比起上億個母體，或許不夠比例，在這裡加入 hard examples。選取 hard negative example 的原則：在 Personalized PageRank 分數在對於 q 的排名 2000-5000 名間，主要找不見得相似於 positive sample 但相似於 q 。



Query



Positive Example



Random Negative



Hard Negative

除此之外，在第 n 個 epoch training 加入 $n-1$ 個 hard negative items 進 negative sampling set。

使用 1.2 billion pair positive example，每個 batch 500 個 negative example，6 個 hard example per pin。

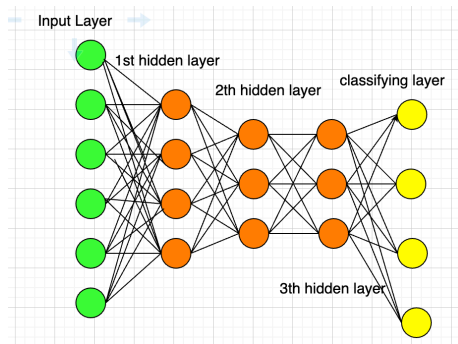
Feature used for learning

$$x_q = [x_q^1, x_q^2, x_q^3]$$

- $x_q^1 \in \mathbb{R}^{4096}$ visual embedding : VGG output
- $x_q^2 \in \mathbb{R}^{256}$ textual embedding : title+description 使用 w2v 的方式訓練結果
- $x_q^3 = \log |E(q)| \in \mathbb{R}^1$

假設 $\text{MLP} : \mathbb{R}^{4393} \rightarrow \mathbb{R}^{128}$ 將每個 item map 到低維空間，context item 指的是 item q 放在同一個 board 裡的其他 item concatenate 在一起，若此 item 被存到許多 board，group by context 會 group 出不同組向量，再 reduce by pooling，輸出單一向量，為此 item q 的 embedding 若要做 board embedding，同樣的執行過程，但輸入為上述的 item embedding(node embedding)。

Multilayer Perceptron



- MLP 是深度模型 (deep neural network, DNN) 的一種
- $\mathbf{h}_{\ell+1} = W\mathbf{h}_{\ell} + b$ or $\mathbf{h}_{\ell+1} = \sigma(W\mathbf{h}_{\ell} + b)$, where $W \in \mathbb{R}^{d_{\ell} \times d_{\ell+1}}$ for $\ell = 1, \dots, L-1$

Baseline comparison

Method	Hit-rate	MRR
Visual	17%	0.23
Annotation	14%	0.19
Combined	27%	0.37
max-pooling	39%	0.37
mean-pooling	41%	0.51
mean-pooling-xent	29%	0.35
mean-pooling-hard	46%	0.56
PinSage	67%	0.59

Visual：單純使用 VGG embedding vector 找歐式距離最近的 itempin)

Annotation：單純使用 title+description vector


Combined：concatenate visual+annotation vector 進 2 層 MLP

hit-rate：推薦 500 個，有被點擊的機率

MRR=Mean Reciprocal Rank
$$MRR = \frac{1}{n} \sum_{(q,i) \in \mathcal{L}} \frac{1}{\lceil R_{i,q}/100 \rceil}$$

Triplet Loss


論文：FaceNet: A Unified Embedding for Face Recognition and Clustering



$$L = \sum_{i=1}^N \max(0, \|f(x_i^a) - f(x_i^p)\|_2^2 - \|f(x_i^a) - f(x_i^n)\|_2^2 + \alpha)$$

$f(x) \in \mathbb{R}^d$: feature vector, x_i^a : anchor image, x_i^p : positive image, x_i^n : negative image

Feature Engineering



萃取 x 特色的

$$f: \mathcal{X} \rightarrow \mathcal{Y}$$

$$x \mapsto f(x)$$

- Language Features 將字變向量
- Image Feature 將照片變相量

Multiple classification

參考論文：Improved deep metric learning with multi-class n-pair loss objective

The $(N + 1)$ -tuple loss is to optimize identifying a positive sample from N negative samples. It is

$$L(x_i, x_j) = \mathbf{1}\{y_i = y_j\} \|f(x_i) - f(x_j)\|^2 + \mathbf{1}\{y_i \neq y_j\} \max(0, \alpha - \|f(x_i) - f(x_j)\|)^2$$

Graph Theory

An undirected simple graph is an ordered pair $G = (V, E)$ comprising

- V is a set of vertices (nodes)
- $E \subseteq \{\{x, y\} | x, y \in V, x \neq y\}$ a set of edges (links or lines) (不允許 loop)

Given a simple graph with n vertices, its Laplacian matrix L

$$L = D - A,$$

where D is the degree matrix and A is the adjacency matrix., which is a binary matrix.

Symmetric normalized Laplacian matrix

The symmetric normalized Laplacian matrix is


$$\mathcal{L} = L^{sym} = D^{-\frac{1}{2}} L D^{-\frac{1}{2}} = I - D^{-\frac{1}{2}} A D^{-\frac{1}{2}}$$

if G has no isolated points.

Example of finding Laplacian

Example

Find the degree matrix D , Adjacency matrix A , Laplacian matrix L , symmetric normalized Laplacian matrix L^{sym}



Example of finding Laplacian

$$A = \begin{pmatrix} 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}, D = \begin{pmatrix} 2 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}. \text{ The Laplacian matrix is}$$

$$L = D - A = \begin{pmatrix} 2 & -1 & 0 & 0 & -1 & 0 \\ -1 & 3 & -1 & 0 & -1 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ -1 & -1 & 0 & -1 & 3 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 \end{pmatrix} \text{ The normalized Laplacian matrix}$$

$$\mathcal{L} = L^{sym} = I - D^{-\frac{1}{2}} A D^{-\frac{1}{2}} = \begin{pmatrix} 1 & -\frac{1}{\sqrt{6}} & 0 & 0 & -\frac{1}{\sqrt{6}} & 0 \\ -\frac{1}{\sqrt{6}} & 1 & -\frac{1}{\sqrt{6}} & 0 & -\frac{1}{3} & 0 \\ 0 & -\frac{1}{\sqrt{6}} & 1 & -\frac{1}{\sqrt{6}} & 0 & 0 \\ 0 & 0 & -\frac{1}{\sqrt{6}} & 1 & -\frac{1}{3} & -\frac{1}{\sqrt{3}} \\ -\frac{1}{\sqrt{6}} & -\frac{1}{3} & 0 & -\frac{1}{3} & 1 & 0 \\ 0 & 0 & 0 & -\frac{1}{\sqrt{3}} & 0 & 1 \end{pmatrix}$$

Laplace operator

Consider a function $f: V \rightarrow \mathbb{R}$ as a vector-valued of function $\mathbf{f}^T = (f(1), \dots, f(n))$ with $|V| = n$.

- The adjacency matrix can be viewed as an operator

$$\mathbf{g} = A\mathbf{f}, \quad g(i) = A(f)(i) = \sum_j A_{ij}f(j)$$

- It can be viewed as a quadratic form

$$\mathbf{f}^T A \mathbf{f} = \sum_{i,j} f(i) A_{ij} f(j) = \sum_{e_{ij}} f(i) f(j)$$

The incidence matrix of a directed graph


Definition

The incidence matrix of a graph is a $|E| \times |V|$ ($m \times n$) matrix defined as follows:

$$\nabla := (\nabla)_{ev} = \begin{cases} -1 & \text{if edge } e \text{ leaves } v \\ 1 & \text{if edge } e \text{ enters } v \\ 0 & \text{otherwise} \end{cases}$$

The incidence matrix: A discrete differential operator

- The mapping $\nabla : \mathbf{f} \mapsto \nabla \mathbf{f}$ is known as the co-boundary mapping of the graph.
- $\nabla \mathbf{f}(e) = v_{in} - v_{out}$



$$\nabla(f) = \begin{pmatrix} -1 & 1 & 0 \\ 0 & -1 & 1 \end{pmatrix} \begin{pmatrix} f(1) \\ f(2) \\ f(3) \end{pmatrix} = \begin{pmatrix} f(2) - f(1) \\ f(3) - f(2) \end{pmatrix}$$

The Laplacian operator of a directed graph

- $L = \nabla^T \nabla$
- $L(\mathbf{f})(i) = \sum_{j=1}^n (\nabla)_{ji} (f(i) - f(j))$
- $L = D - A$ in matrix form

The Laplacian operator of an undirected graph

- The Laplacian as an operator:

$$L(\mathbf{f})(i) = \sum_{j=1}^n A_{ij}(f(i) - f(j))$$

- As a quadratic form

$$\begin{aligned}\mathbf{f}^T L \mathbf{f} &= \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n A_{ij} (f(i) - f(j))^2 = (f, L(\mathbf{f})) \\ &= (\nabla \mathbf{f}, \nabla \mathbf{f}) = \|\nabla f\|_A^2 = (-\Delta f, f)\end{aligned}$$

- L is symmetric and positive semi-definite.
- L has n non-negative, real-valued eigenvalues: $0 = \lambda_1 \leq \dots \leq \lambda_n$

Graph Fourier Transform

Definition

A graph signal $f: V \rightarrow \mathbb{R}$ is a function on the vertices of the graph G . Let λ_i and u_i be the i -th eigenvalue and eigenvector of the Laplacian matrix L . The graph Fourier transform \hat{f} :

$$\mathcal{GF}[f](\lambda_i) = \hat{f}(\lambda_i) = \langle f, u_i \rangle = \sum_{k=1}^n u_i(k) f(k),$$

where $\{u_i\}_{i=1}^n$ forms an orthonormal basis of \mathbb{R}^n . Thus

$$\mathcal{GF}[\mathbf{f}] = \hat{\mathbf{f}} = U^T \mathbf{f}$$

Recall: $\mathbf{f} = \sum_{i=1}^n \langle \mathbf{f}, u_i \rangle u_i = \sum_{i=1}^n \hat{\mathbf{f}}(i) u_i$

The inverse graph Fourier transform

Definition


The inverse graph Fourier transform is

$$\mathcal{IGF}[\hat{f}](i) = f(i) = \sum_{k=1}^{n-1} \hat{f}(\lambda_k) u_k(i).$$

In matrix form

$$\mathcal{IGF}[\hat{\mathbf{f}}] = \mathbf{f} = U\hat{\mathbf{f}}$$

Example



$L = D - A = \begin{pmatrix} 2 & -1 & -1 \\ -1 & 1 & 0 \\ -1 & 0 & 1 \end{pmatrix}$, The characteristic polynomial of L is

$p(\lambda) = \det(L - \lambda I) = -\lambda(\lambda - 1)(\lambda - 3)$. The unit corresponding

eigenvector are $u_{\lambda_1} = \begin{pmatrix} \frac{1}{\sqrt{3}} \\ \frac{1}{\sqrt{3}} \\ \frac{1}{\sqrt{3}} \end{pmatrix}$, $u_{\lambda_2} = \begin{pmatrix} 0 \\ \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{pmatrix}$, $u_{\lambda_3} = \begin{pmatrix} \frac{-2}{\sqrt{6}} \\ \frac{1}{\sqrt{6}} \\ \frac{1}{\sqrt{6}} \end{pmatrix}$.

$U = (u_{\lambda_1} \quad u_{\lambda_2} \quad u_{\lambda_3})$. Then $L = U \cdot \text{diag}(\lambda_1, \lambda_2, \lambda_3) U^T$

Example of Graph Fourier Transform

$U = \begin{pmatrix} \frac{1}{\sqrt{3}} & 0 & \frac{-2}{\sqrt{6}} \\ \frac{1}{\sqrt{3}} & -\frac{1}{\sqrt{2}} & \frac{1}{\sqrt{6}} \\ \frac{1}{\sqrt{3}} & \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{6}} \end{pmatrix}$. If $\mathbf{f} = \begin{pmatrix} f(1) \\ f(2) \\ f(3) \end{pmatrix} = \begin{pmatrix} 1 \\ -1 \\ 1 \end{pmatrix}$. Then the graph Fourier transform of \mathbf{f} is

$$\mathcal{GF}[\mathbf{f}] = \hat{\mathbf{f}} = U^T \mathbf{f} = \begin{pmatrix} \frac{1}{\sqrt{3}} & \frac{1}{\sqrt{3}} & \frac{1}{\sqrt{3}} \\ 0 & -\frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \\ \frac{-2}{\sqrt{6}} & \frac{1}{\sqrt{6}} & \frac{1}{\sqrt{6}} \end{pmatrix} \begin{pmatrix} 1 \\ -1 \\ 1 \end{pmatrix} = \begin{pmatrix} \frac{1}{\sqrt{3}} \\ \frac{2}{\sqrt{2}} \\ \frac{-2}{\sqrt{6}} \end{pmatrix}$$

The inverse graph Fourier transform of $\hat{\mathbf{f}}$ is

$$\mathcal{IGF}[\hat{\mathbf{f}}] = \mathbf{f} = U \hat{\mathbf{f}} = \begin{pmatrix} \frac{1}{\sqrt{3}} & 0 & \frac{-2}{\sqrt{6}} \\ \frac{1}{\sqrt{3}} & -\frac{1}{\sqrt{2}} & \frac{1}{\sqrt{6}} \\ \frac{1}{\sqrt{3}} & \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{6}} \end{pmatrix} \begin{pmatrix} \frac{1}{\sqrt{3}} \\ \frac{2}{\sqrt{2}} \\ \frac{-2}{\sqrt{6}} \end{pmatrix} = \begin{pmatrix} 1 \\ -1 \\ 1 \end{pmatrix}$$

Graph convolution

Then $\mathcal{GF}[\mathbf{f}] = \hat{\mathbf{f}} = U^T \mathbf{f}$. In particular, for graph convolution

$$(\mathbf{f} * \mathbf{g}) = \mathcal{IGF}[\mathcal{GF}[(\mathbf{f} * \mathbf{g})]]$$

That is,

$$(\mathbf{f} * \mathbf{g})(i) = \sum_{k=1}^n \hat{\mathbf{f}}(\lambda_k) \hat{\mathbf{g}}(\lambda_k) u_k(i)$$
$$(\mathbf{f} * \mathbf{g})_G = U \begin{pmatrix} \hat{\mathbf{g}}(\lambda_1) & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \hat{\mathbf{g}}(\lambda_n) \end{pmatrix} U^T \mathbf{f}$$

參考論文：SEMI-SUPERVISED CLASSIFICATION WITH GRAPH CONVOLUTIONAL NETWORKS 2017

假設任務是要分類 graph 裡的 node，但只有部分答案 (some labels are available) · 這個監督問題可以被圖的結構弭平

$$\mathcal{L} = \mathcal{L}_0 + \lambda \cdot \mathcal{L}_{reg},$$

where \mathcal{L}_0 是監督式有標籤的 loss、

$$\mathcal{L}_{reg} = \frac{1}{2} \sum_{i,j} A_{ij} \|f(X_i) - f(X_j)\|^2 = f(X)^T \Delta f(X) = \frac{1}{2} \|\nabla f\|_A^2$$

where $\Delta = D - A$ the unnormalized graph Laplacian of an undirected graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ with N nodes $v_i \in \mathcal{V}$, edges $(v_i, v_j) \in \mathcal{E}$, an adjacency matrix $A \in \mathcal{R}^{N \times N}$ (binary or weighted) and a degree matrix $D_{ii} = \sum_{j=1}^N A_{ij}$ · 此 loss function 是基於一個聯通的圖會分享同一個 label，這個假設有可能會限制 model 的容納量

Fast Approximate Convolution on Graph

建立一個以圖為基礎的多層圖卷機 (multi-layer Graph Convolutional Network: GCN)

$$H^{(\ell+1)} = \sigma \left(\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} H^{(\ell)} W^{(\ell)} \right)$$

where $\tilde{A} = A + I_N$: the adjacency matrix of the undirected graph \mathcal{G} with added self-connection,

$$\tilde{D}_{ii} = \sum_{j=1}^N \tilde{A}_{ij}$$

$W^{(\ell)} \in \mathbb{R}^{D \times D}$: 第 ℓ 層的權重

σ : activation function 可以是 ReLU 或其他

$H^{(\ell)} \in \mathbb{R}^{N \times D}$: output of the ℓ -th hidden layer

$H^{(0)} = X$

Spectral Graph Convolution

Given a signal $x \in \mathbb{R}^N$, 定義 spectral graph convolution with a filter:
 $g_\theta = \text{diag}(\theta)$, where $\theta \in \mathbb{R}^N$

$$g_\theta \star x = U g_\theta U^T x = U \begin{pmatrix} \theta_1 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \theta_N \end{pmatrix} U^T x$$

where U is the matrix of eigenvectors of the symmetric normalized graph Laplacian \mathcal{L}

$$\begin{aligned} \mathcal{L} &= D^{-\frac{1}{2}}(\Delta)D^{-\frac{1}{2}} = D^{-\frac{1}{2}}(D - A)D^{-\frac{1}{2}} \\ &= I_N - D^{-\frac{1}{2}}AD^{-\frac{1}{2}} = U\Lambda U^T, \end{aligned}$$

Λ 是 \mathcal{L} 的對角化矩陣（特徵值對角矩陣， U 的 column vector \mathcal{L} 的特徵向量

Approximation of Chebyshev polynomials

$g_{\theta}(\Lambda)$ is a function of eigenvalues of \mathcal{L} · 但計算成本太大，這可被 truncated expansion in terms of Chebyshev polynomial $T_k(x)$ up to K-th order

$$g_{\theta'}(\Lambda) \approx \sum_{k=0}^K \theta'_k T_k(\tilde{\Lambda})$$

with rescale $\tilde{\Lambda} = \frac{2}{\lambda_{\max}}\Lambda - I_N$, λ_{\max} denote the largest eigenvalue of \mathcal{L} .
(因為 T_k 定義範圍在 $[-1,1]$ 之間，所以需要 rescale 輸入) $\theta' \in \mathbb{R}^K$ is a vector of Chebyshev coefficients $T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x)$ with $T_0(x) = 1$ and $T_1(x) = x$
重新定義 a convolution of a signal x with filter $g_{\theta'}$

$$g_{\theta'} \star x \approx \sum_{k=0}^K \theta'_k T_k(\tilde{L})x,$$

where $\tilde{L} = \frac{2}{\lambda_{\max}}L - I_N$ · 因為 $(U\Lambda U^T)^k = U\Lambda^k U^T$

Layer-Wise Linear Model

$$K = 1, \lambda_{\max} \approx 2$$

$$g_{\theta'} \star x \approx \theta'_0 x + \theta'_1 (L - I_N) x = \theta'_0 - \theta'_1 D^{-\frac{1}{2}} A D^{-\frac{1}{2}} x$$

減低參數 $\theta = \theta'_0 = -\theta'_1$ ，上式變成

$$g_{\theta} \star x \approx \theta (I_N + D^{-\frac{1}{2}} A D^{-\frac{1}{2}}) x$$

renormalization trick: $I_N + D^{-\frac{1}{2}} A D^{-\frac{1}{2}} \rightarrow \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$ with $\tilde{A} = A + I_N$ and $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$

若推廣這個定義到 a signal $X \in \mathbb{R}^{N \times C}$ with C input channels and F filters ($\Theta \in \mathbb{R}^{C \times F}$ matrix of filter parameters)

$$Z = \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} X \Theta \in \mathbb{R}^{N \times F}$$

is the convolved signal

Semi-Supervised Node classification


考慮 two-layer GCN for semi-supervised node classification: let $\hat{A} = \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$

$$Z = f(X, A) = \text{softmax}(\hat{A}(\text{ReLU}(\hat{A}XW^{(0)}))W^{(1)})$$


監督式部分的 loss : cross-entropy

$$\mathcal{L}_0 = - \sum_{\ell \in Y_L} \sum_{f=1}^F Y_{\ell f} \log Z_{\ell f}$$

where Y_L is the set of node indices that have labels .



(a) Graph Convolutional Network



(b) Hidden layer activations

論文：Graph Convolutional Matrix Completion 2017

基於 user-item 的 Bipartite graph，利用 Graph Auto-Encoder 來預測使用者的評分


Let a rating matrix $M \in \mathbb{N}^{N_u \times N_v}$, where N_u is the number is users and N_v is the number of items

常用方式：協同過濾

$$M_{ij} = \begin{cases} r & \text{if user } i \text{ rated item } j \\ 0 & \text{otherwise} \end{cases}$$


for $r \in \{1, \dots, R\}$

整體模型



- Bipartite graph $G = (\mathcal{W}, \mathcal{E}, \mathcal{R})$, where $\mathcal{U} = \{u_i\}_{i=1}^{N_u}$: users nodes, $\mathcal{V} = \{v_j\}_{j=1}^{N_v}$: item nodes, $\mathcal{W} = \mathcal{U} \cap \mathcal{V}$, edge (u_i, r, v_j) for $r \in \{1, \dots, R\} = \mathcal{R}$
- Graph encoders: $Z = f(X, A)$, where $X \in \mathbb{R}^{N \times D}$ for $N = N_u + N_v$ 所有 nodes, A : adjacency matrix
output: $Z = [z_1^T, \dots, z_N^T]^T \in \mathbb{R}^{N \times E}$
- decoder $g(Z) = \check{A}$, $g(z_i, z_j) = \check{A}_{ij}$

Graph auto-encoders



- Encoder $f(X, M_1, \dots, M_R) = [U, V]$, where $M_r \in \{0, 1\}^{N_u \times N_v}$, for $r \in \mathcal{R}$, $U \in \mathbb{R}^{N_u \times E}$, $V \in \mathbb{R}^{N_v \times E}$
- Decoder $\tilde{M} = g(U, V)$ rating matrix of shape $N_u \times N_v$

Graph convolution encoder for users

- The hidden vector for each user i

$$h_i = \sigma \left[\sum_{j \in \mathcal{N}_{i,1}} \mu_{j \rightarrow i,1}, \dots, \sum_{j \in \mathcal{N}_{i,R}} \mu_{j \rightarrow i,R} \right],$$

where $\mu_{j \rightarrow i,r} = \frac{1}{c_{ij}} W_r x_j$: edge-type rating result from item j to user i ,
 c_{ij} is a normalization constant $c_{ij} = |\mathcal{N}_i|$ or $c_{ij} = \sqrt{|\mathcal{N}_i| |\mathcal{N}_j|}$ with \mathcal{N}_i
denoting the set of neighbors of node i , and
 $\sigma(x) = \text{ReLU}(x) = \max(0, x)$.

- $\sum_{j \in \mathcal{N}_{i,r}} \mu_{j \rightarrow i,r} = \sum_{j=1}^{N_v} M_r \mu_{j \rightarrow i,r}$
- The output of GAE for the user vector $u_i = \sigma(Wh_i)$

Graph convolution encoder for users

Similarly, item vector 也可透過同樣的方式得到：

•

$$\mu_{i \rightarrow j, r} = \frac{1}{c_{ij}} W_r x_i$$

•

$$h_j = \sigma \left[\sum_{i \in N_{j,1}} \mu_{i \rightarrow j,1}, \dots, \sum_{i \in N_{j,R}} \mu_{i \rightarrow j,R} \right]$$

- The output of GAE for the item vector $v_j = \sigma(Wh_j)$

Decoder

再來經過一個 bilinear decoder $\check{M} = g(U, V)$:

- The predicted rating

$$\check{M}_{ij} = g(u_i, v_j) = \mathbb{E}_{p(\check{M}_{ij}=r)}[r] = \sum_{r \in R} r p(\check{M}_{ij} = r)$$

- The decoder produces a probability distribution of the rating between user i and item j with M_{ij} ,

$$p(\check{M}_{ij} = r) = \frac{e^{u_i^T Q_r v_j}}{\sum_{s \in R} e^{u_i^T Q_s v_j}},$$

$Q_r \in \mathbb{R}^{E \times E}$ 是學習參數。

Training

Loss function:

$$L = - \sum_{i,j \in \Omega: \Omega_{ij}=1} \sum_{r=1}^R \mathbf{1}[r = M_{ij}] \log p(\check{M}_{ij} = r),$$

where $\Omega \in \{0, 1\}^{N_u \times N_v}$ serves as a mask for unobserved ratings, such that ones occur for elements corresponding to observed ratings in M , and zeros for unobserved ratings.

Matrix representation



$$Z = \begin{bmatrix} U \\ V \end{bmatrix} = f(X, M_1, \dots, M_R) = \sigma \left(\begin{bmatrix} H_u \\ H_v \end{bmatrix} W^T \right),$$

where $X = \begin{bmatrix} X_u \\ X_v \end{bmatrix}$, $\begin{bmatrix} H_u \\ H_v \end{bmatrix} = \sigma \left(\sum_{r=1}^R D^{-1} \mathcal{M}_r X W_r^T \right)$,

$$\mathcal{M}_r = \begin{pmatrix} 0 & M_r \\ M_r^T & 0 \end{pmatrix}, U \in \mathbb{R}^{N_u \times E}, V \in \mathbb{R}^{N_v \times E}.$$

- Feature representation

為了區分 user node 與 item node，可加入 user side information 與 item side information：

user: $u_i = \sigma(W h_i + W_2^f f_i)$, $f_i = \sigma(W_1^f x_i^f + b)$,

item: $v_j = \sigma(W h_j + W_2^f f_j)$, $f_j = \sigma(W_1^f x_j^f + b)$

在 GCN 輸入時的 $X = [X_u, X_v]$ 使用 one-hot。所以 X 是一個 identity matrix。

Training

- Mini-batch
- 每次只採固定數量的 user-item 對
- Node Dropout
為了使模型泛化到未觀測到的評分，在訓練中使用 dropout，以一定機率 p_{dropout} 刪除特定節點所有的傳出值。
- Weight sharing