

人工智慧專題

劉瓊如

October 14, 2022

1 Natural Language Processing (NLP)


- Language Models
- Gradient estimation

2 Classification

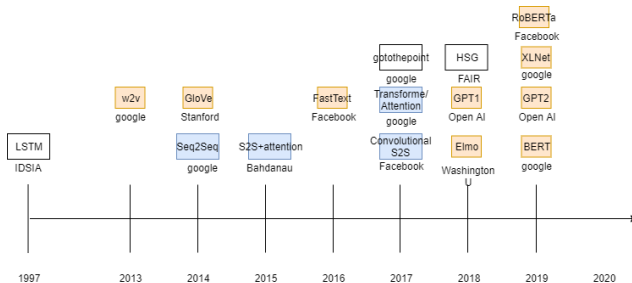
- Metrics

Language Models



- word2v
 - ▶ 優點 (vector similarity)
 - ▶ 缺點新字要重新訓練
- GloVe 字頻
- ELMo (next word)
- BERT (克漏字填空)
 - ▶ 優點：character level，可以組新字，組 sentence vector
 - ▶ 缺點：wors similarity 較差，例如有重複字："雪白"、"亮白" word similarity 接近，"白目" 字義不接近，但 word similarity 接近
- GPT: 適合生成模型
- XLNet
- Roberta



NLP History



Recurrent Neural Networks (RNNs)



The repeating module in a standard RNN contains a single layer.

$$h_t(x_t) = \sigma(h_{t-1}, x_t) = \sigma(W_{hh}h_{t-1} + W_{hx}x_t + b)$$

$$\hat{y}_t(x_t) = W_{hy}h_t(x_t) + b_y$$

Next word prediction

$$\mathbf{p}_t = \text{softmax}(\hat{\mathbf{y}}) \in \mathbb{R}^{|V|}$$

$$L = -\frac{1}{T} \sum_{t=1}^T \mathbf{y}_t \log \mathbf{p}_t$$

\mathbf{p}_t : output probability distribution of the next word

\mathbf{h}_T : sentence representation

$$\frac{\partial L_t}{\partial W} = \sum_{k=1}^t \frac{\partial L_t}{\partial \mathbf{p}_t} \frac{\partial \mathbf{p}_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_k} \frac{\partial \mathbf{h}_k}{\partial W}$$

$$\frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_k} = \Pi_{j=k+1}^t \frac{\partial \mathbf{h}_j}{\partial \mathbf{h}_{j-1}} = \Pi_{j=k+1}^t \text{diag}(\sigma'(W\mathbf{h}_{j-1} + \dots)) \times W$$

Gradient estimation

$$\frac{\partial L}{\partial W} = \sum_{t=1}^T \sum_{k=1}^t \frac{\partial L_t}{\partial \mathbf{p}_t} \frac{\partial \mathbf{p}_t}{\partial \mathbf{h}_t} (\Pi_{j=k+1}^t \frac{\partial \mathbf{h}_j}{\partial \mathbf{h}_{j-1}}) \frac{\partial \mathbf{h}_k}{\partial W}$$

If $\|\frac{\partial \mathbf{h}_j}{\partial \mathbf{h}_{j-1}}\| \leq \|\text{diag}(\sigma'(Wh_{j-1}))\| \|W\| \leq \beta_h \beta_W$ by taking their L_2 -norm


$$\|A\|_2 = \sup_{\mathbf{x} \neq 0} \frac{\|A\mathbf{x}\|_2}{\|\mathbf{x}\|_2} = \sigma_{\max}(A)$$

$$\|\frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_k}\| \leq \|\Pi_{j=k+1}^t \frac{\partial \mathbf{h}_j}{\partial \mathbf{h}_{j-1}}\| \leq (\beta_h \beta_W)^{t-k}$$

is easily large or small if $t - k$ is sufficiently large. This could cause the Gradient Explosion Problem or Gradient Vanishing Problem

LSTM

論文：Long Short-Term Memory in Recurrent Neural Networks (Hochreiter and Schmidhuber)1997



$$i_t = \sigma(W_x^{(i)} x_t + W_h^{(i)} h_{t-1} + b_i) \quad \text{input gate}$$

$$f_t = \sigma(W_x^{(f)} x_t + W_h^{(f)} h_{t-1} + b_f) \quad \text{forget gate}$$


$$o_t = \sigma(W_x^{(o)} x_t + W_h^{(o)} h_{t-1} + b_o) \quad \text{output gate}$$

$$\tilde{c}_t = \tanh(W_x^{(\tilde{c})} x_t + W_h^{(\tilde{c})} h_{t-1} + b_{\tilde{c}}) \quad \text{temporary memory}$$

$$c_t = i_t \odot \tilde{c}_t + f_t \odot c_{t-1} \quad \text{memory gate}$$

$$h_t = o_t \odot \tanh(c_t),$$

Backpropagation



$$\frac{\partial L_t}{\partial W_{(\cdot)}} = \begin{cases} \sum_{k=1}^t \frac{\partial L_t}{\partial \mathbf{p}_t} \frac{\partial \mathbf{p}_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{c}_t} \frac{\partial \mathbf{c}_t}{\partial \mathbf{c}_k} \frac{\partial \mathbf{c}_k}{\partial (\cdot)_k} \frac{\partial (\cdot)_k}{\partial W_{(\cdot)}} & (\cdot) = i, f, \tilde{c} \\ \sum_{k=1}^t \frac{\partial L_t}{\partial \mathbf{p}_t} \frac{\partial \mathbf{p}_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_k} \frac{\partial \mathbf{h}_k}{\partial \mathbf{o}_k} \frac{\partial \mathbf{o}_k}{\partial W^{(\circ)}} & \text{otherwise} \end{cases}$$

$$\frac{\partial \mathbf{c}_t}{\partial \mathbf{c}_k} = \prod_{s=k+1}^t \frac{\partial \mathbf{c}_s}{\partial \mathbf{c}_{s-1}}$$

$$\frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_k} = \prod_{s=k+1}^t \frac{\partial \mathbf{h}_s}{\partial \mathbf{h}_{s-1}} = \prod_{s=k+1}^t \frac{\partial \mathbf{h}_s}{\partial \mathbf{c}_s} \frac{\partial \mathbf{c}_s}{\partial \mathbf{h}_{s-1}}$$

$$\begin{aligned}
\frac{\partial \mathbf{c}_t}{\partial \mathbf{c}_{t-1}} &= \frac{\partial}{\partial \mathbf{c}_{t-1}} [\mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{c}}_t] \\
&= \text{diag}(\mathbf{f}_t) + \text{diag}(\mathbf{c}_{t-1}) \frac{\partial \mathbf{f}_t}{\partial \mathbf{c}_{t-1}} + \text{diag}(\tilde{\mathbf{c}}_t) \frac{\partial \mathbf{i}_t}{\partial \mathbf{c}_{t-1}} + \text{diag}(\mathbf{i}_t) \frac{\partial \tilde{\mathbf{c}}_t}{\partial \mathbf{c}_{t-1}} \\
&= A_t + B_t + C_t + D_t
\end{aligned}$$

$$\frac{\partial \mathbf{f}_t}{\partial \mathbf{c}_{t-1}} = \frac{\partial \mathbf{f}_t}{\partial \mathbf{h}_{t-1}} \frac{\partial \mathbf{h}_{t-1}}{\partial \mathbf{c}_{t-1}} = W_h^{(f)} \text{diag}(\sigma'(W_h^{(f)} h_{t-1})) \text{diag}(\mathbf{o}_t) \text{diag}(\tanh'(c_{t-1}))$$

$$\frac{\partial \mathbf{i}_t}{\partial \mathbf{c}_{t-1}} = \frac{\partial \mathbf{i}_t}{\partial \mathbf{h}_{t-1}} \frac{\partial \mathbf{h}_{t-1}}{\partial \mathbf{c}_{t-1}} = W_h^{(i)} \text{diag}(\sigma'(W_h^{(i)} h_{t-1})) \text{diag}(\mathbf{o}_t) \text{diag}(\tanh'(c_{t-1}))$$


$$\frac{\partial \tilde{\mathbf{c}}_t}{\partial \mathbf{c}_{t-1}} = \frac{\partial \tilde{\mathbf{c}}_t}{\partial \mathbf{h}_{t-1}} \frac{\partial \mathbf{h}_{t-1}}{\partial \mathbf{c}_{t-1}} = W_h^{(\tilde{c})} \text{diag}(\tanh'(W_h^{(\tilde{c})} h_{t-1})) \text{diag}(\mathbf{o}_t) \text{diag}(\tanh'(c_{t-1}))$$

Preventing the error gradients from vanishing

$$\frac{\partial L_t}{\partial W_{(\cdot)}} = \sum_{k=1}^t \frac{\partial L_t}{\partial \mathbf{p}_t} \frac{\partial \mathbf{p}_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{c}_t} (\Pi_{t=k+1}^t [A_t + B_t + C_t + D_t]) \frac{\partial \mathbf{c}_k}{\partial W_{(\cdot)}}$$

因為梯度包含 forget gates 的激活向量，可以控制梯度的大小（ A_t 項乘的次方較少，值較大）。可以透過參數更新，讓神經元選擇遺忘或保留記憶。

BLSTM




ELMo

論文：Deep contextualized word representations

(Allen Institute for Artificial Intelligence+University of Washington)

ELMo=(Embeddings from **L**anguage **M**odels)



ELMo

Given a sentence of tokens (t_1, t_2, \dots, t_N)

$$L = \sum_{k=1}^N \left(\log p(t_k | t_1, \dots, t_{k-1}; \vec{\Theta}_{\text{LSTM}}) + \log p(t_k | t_{k+1}, \dots, t_N; \vec{\Theta}_{\text{LSTM}}) \right),$$

For each token t_k , a L -layer biLM computes a set of $2L + 1$ representations

$$\begin{aligned} R_k &= \{x_k^{LM}, \vec{h}_{k,j}^{LM}, \overleftarrow{h}_{k,j}^{LM} | j = 1, \dots, L\} \\ &= \{h_{k,j}^{LM}, j = 0, \dots, L\}, \end{aligned}$$

where $h_{k,0}^{LM}$ is the token layer and $h_{k,j}^{LM} = [\vec{h}_{k,j}^{LM}; \overleftarrow{h}_{k,j}^{LM}]$ for each biLM layer. For a downstream task, each word t_k has


$$\mathbf{ELMo}_k^{\text{task}} = E(R_k; \Theta^{\text{task}}) = \gamma^{\text{task}} \sum_{j=0}^L s_j^{\text{task}} h_{k,j}^{LM}$$

for different tasks.

BERT

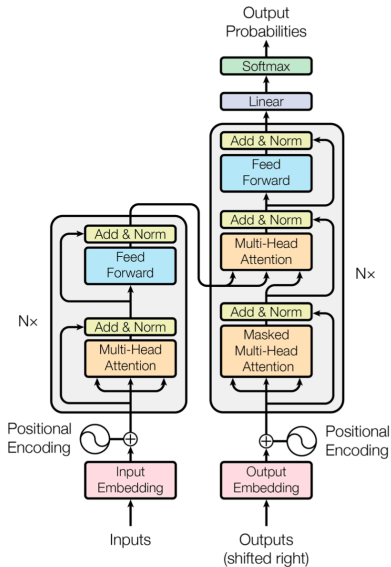
論文: Pre-training of Deep Bidirectional Transformers for Language Understanding

BERT=**B**idirection **E**ncoder **R**epresentation from **T**ransformers 縮寫



- Encoder of Transformer
- Self-supervised Learning
- Cloze

Transformer



Transformer

- 捨棄 RNN 的序列式 input
- **Positional Encoding**: 由於一個句子可以同時進模型，為了可以區分順序

$$\text{PE}(\text{pos}, 2i) = \sin\left(\frac{\text{pos}}{10000^{\frac{2i}{d_{\text{model}}}}}\right)$$

$$\text{PE}(\text{pos}, 2i + 1) = \cos\left(\frac{\text{pos}}{10000^{2i/d_{\text{model}}}}\right)$$

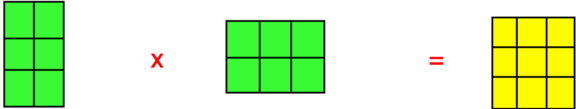
pos 代表的是位置，i 代表的是維度，偶數位置的文字會透過 sin 函數進行轉換，奇數位置的文字則透過 cos 函數進行轉換

- **Attention Mechanism**


$$\text{attention}(Q, K, V) = \text{softmax}\left(\frac{Q \cdot K^T}{\sqrt{d_k}}\right) \cdot V,$$

$$Q = XW_Q, K = XW_K, V = XW_V, X \in \mathbb{R}^{n \times d}, W_Q, W_K \in \mathbb{R}^{d \times d_k}, \\ W_V \in \mathbb{R}^{d \times d_v}$$

Attention Score

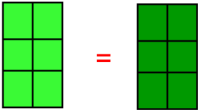
$$Q \times K^T =$$


A diagram illustrating the multiplication of two matrices. On the left is a green 3x2 matrix labeled Q . In the middle is a green 3x3 matrix labeled K^T . To the right of the green matrices is a red 'x' symbol, followed by an equals sign. To the right of the equals sign is a yellow 3x3 matrix, representing the result of the matrix multiplication.

$$\text{softmax} =$$


A diagram showing a 3x3 yellow matrix followed by an equals sign and a 3x3 matrix with numerical values. The numerical matrix has a header row with x_1 , x_2 , and x_3 in red. The values are: x_1 row: 0.6, 0.3, 0.1; x_2 row: 0.2, 0.5, 0.3; x_3 row: 0.1, 0.2, 0.7.

	x_1	x_2	x_3
x_1	0.6	0.3	0.1
x_2	0.2	0.5	0.3
x_3	0.1	0.2	0.7

$$\text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) \times V =$$


A diagram showing the multiplication of a 3x2 green matrix by a 3x2 dark green matrix labeled V . The green matrix is preceded by the expression $\text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)$. To the right of the green matrix is a red 'x' symbol, followed by an equals sign. To the right of the equals sign is a dark green 3x2 matrix, representing the result of the matrix multiplication.


Self-Attention Sublayer

- Multi-head attention:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O$$

$$\text{head}_i = \text{attention}(QW_i^Q, KW_i^K, VW_i^V)$$

$W_i^Q, W_i^K \in \mathbb{R}^{d_{\text{model}} \times d_k}$, $W_i^V \in \mathbb{R}^{d_{\text{model}} \times d_v}$. Default $h = 8$,
 $d_k = d_v = d_{\text{model}} = 64$.



Feed-Forward Sublayer

- **Position-wise Feed-Forward Networks:**

$$\text{FFN}(\mathbf{x}) = \max\{0, \mathbf{x}W_1 + \mathbf{b}_1\} \cdot W_2 + \mathbf{b}_2,$$


$$W_1 \in \mathbb{R}^{d_{\text{model}} \times d_{\text{ff}}}, W_2 \in \mathbb{R}^{d_{\text{ff}} \times d_{\text{model}}}, d_{\text{ff}} = 4d_{\text{model}}$$

- **Residual Connections:** We employ a residual connection around each of the two sub-layers, followed by layer normalization:

$$\text{LayerNorm}(\mathbf{x} + \text{Sublayer}(\mathbf{x}))$$

防止 gradient vanishing

Layer Normalization




$$(\text{LayerNorm}(\mathbf{x}))_t^i = \frac{(\mathbf{x})_t^i - \mu_i}{\sigma_i}$$

$$\mu_i = \frac{1}{n} \sum_{t=1}^n (\mathbf{x})_t^i$$

$$\sigma_i = \sqrt{\frac{1}{n} \sum_{t=1}^n ((\mathbf{x})_t^i - \mu_i)^2}$$

Embedding

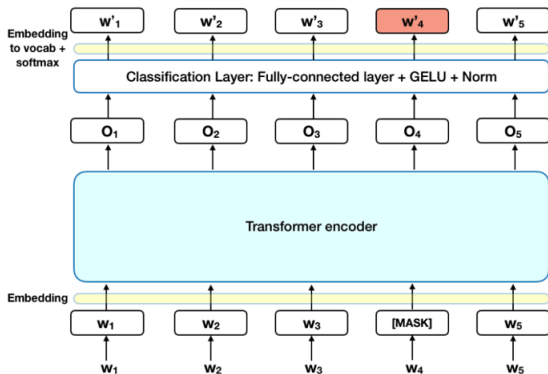
- Input embeddings: 三種 embedding 加在一起
Token embedding + Segment embedding + Position embedding = word embedding + sentence embedding + Position embedding



假定句字最長 512 字 · 不足補 0

Mask of BERT

使用 mask 推測當前的字。



$$L = -\log(p(w_t = \text{Mask}) | w_1, \dots, w_{t-1}, w_{t+1}, \dots, w_L)$$

Mask

克漏字原理

In all of our experiments, we mask 15% of all WordPiece tokens in each sequence at random. Rather than always replacing the chosen words with [MASK], the data generator will do the following:

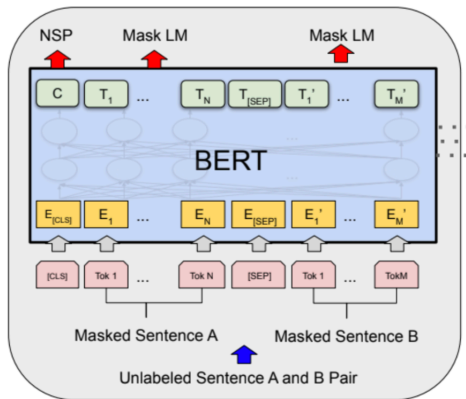
- ① 80% of the time: Replace the word with the [MASK] token, e.g., my dog is hairy → my dog is [MASK]
- ② 10% of the time: Replace the word with a random word, e.g., my dog is hairy → my dog is apple
- ③ 10% of the time: Keep the word unchanged, e.g., my dog is hairy → my dog is hairy. The purpose of this is to bias the representation towards the actual observed word.

雖然只 15% 的字需要預測，不知道哪個字被換掉，所以每個字都要關注。

Pre-training BERT: Next Sentence Prediction

- 判斷是否是下句，分割成對句，若全文有 40 句，拆成 20 對:
pre-train a binarized next-sentence prediction.
Note: Use a document-level corpus rather than a shuffled sentence-level corpus
每兩句話做一個 input，label output: IsNext, NotNext
訓練的時候，將 50% 的句子按原文前後順序放，50% 隨便亂放
- Special token: [CLS], [Mask], [Sep]
- 假設 $A \rightarrow B$ ，若輸入 [CLS] A [Sep] B，則 [CLS] token 需猜測 [IsNext]
- 若輸入 [CLS] B [Sep] A，則 [CLS] token 需猜測 [NotNext]

Pre-training



Pre-training

Training

Task1 的 loss function: 假設 document 有 $2m$ 句, $s_i = \{w_1, \dots, w_{N_i}\}$, N = 總字數:

$$\mathcal{L}_1 = \frac{1}{N} \sum_{i=1}^m \sum_{k=1}^N \log p(w_k = \text{mask} | w_1, \dots, w_{k-1}, w_{k+1}, \dots, w_{N_i}; \Theta),$$

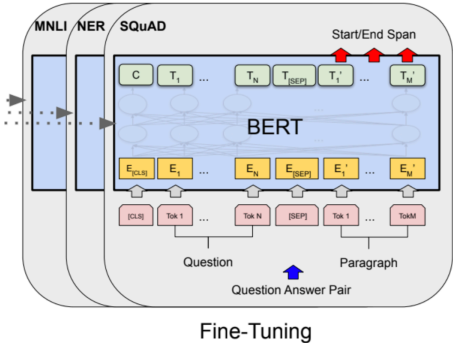
Task2 的 loss function: (假設分成 m 對)
document = $\{s_{1a}, s_{1b}, \dots, s_{ma}, s_{mb}\}$:

$$\mathcal{L}_2 = \frac{1}{m} \sum_{k=1}^m \log p(y | s_{ka}, s_{kb})$$

The model is trained by the total loss

$$\mathcal{L} = \mathcal{L}_1 + \mathcal{L}_2$$

Down Stream Task



Down Stream Task

- Parameters transfer+ classifying layer
- Task: Text classification: NER, POS
- Task: Sentence Prediction: QA, NLI

A: question, B: passage, S: start token embedding, E: end token


- ▶ 第一個字： $\operatorname{argmax}_i P_i$, where $P_i = \frac{e^{S \cdot T_i}}{\sum_j e^{S \cdot T_j}}$: the probability of word i being the start of the answer span.
- ▶ 接下來依序下個字的計算方式：The score of a candidate span from position i to position j is defined as $S \cdot T_i + E \cdot T_j$, and the maximum scoring span where $j \geq i$ is used as a prediction.
- ▶ 最後一個字： $P_i = \frac{e^{E \cdot T_i}}{\sum_j e^{E \cdot T_j}}$ 最高的那個
- The training objective is the sum of the log-likelihoods of the correct start and end positions. (中間不算嗎?)

CoreNLP

Stanford CoreNLP

GPT

論文：Improving Language Understanding by Generative Pre-Training
使用 Transformer 的 decoder



Given an unsupervised corpus of tokens $\mathcal{U} = \{u_1, \dots, u_n\}$

- Architecture:

$$h_0 = UW_e + W_p$$

$$h_\ell = \text{transformer_block}(h_{\ell-1}) \quad \forall i \in [1, n]$$

$$p(x) = \text{softmax}(h_n W_e^T)$$

where $U = (u_{-k}, \dots, u_{-1})$ is the context vector of tokens, n is the number of layers, W_e is the token embedding matrix and W_p is the position embedding matrix.


- Minimizing

$$L_1(\mathcal{U}) = -\log p(u_i | u_{i-k}, \dots, u_{i-1}; \Theta),$$



where k is the size of the context window, and the conditional probability p is modeled using a neural network with parameters Θ .

Masked Attention

softmax




=



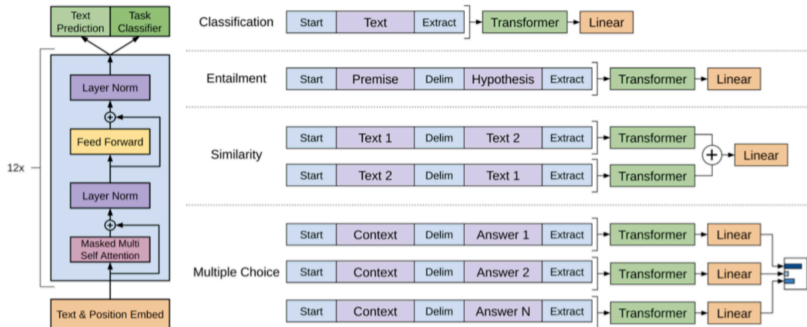
$$\begin{matrix} & V \\ \times & \begin{matrix} v_1 \\ v_2 \\ v_3 \end{matrix} \end{matrix}$$

=

$$\begin{matrix} v_1 \\ 0.2v_1 + 0.8v_2 \\ 0.2v_1 + 0.3v_2 + 0.5v_3 \end{matrix}$$



GPT Task




OpenAI playground

OpenAI

Classification

- Text Classification
- Image Classification



- Model output feature $f(\mathbf{x}) = \mathbf{z} \in \mathbb{R}^d$
- Classifying Layer $\hat{\mathbf{y}} = W\mathbf{z} + \mathbf{b}$, $\mathbf{p} = \text{softmax}(\hat{\mathbf{y}})$, $W \in \mathbb{R}^{C \times d}$
or $\mathbf{p} = \sigma(\hat{\mathbf{y}}) = (\sigma(\hat{y}_1), \dots, \sigma(\hat{y}_C))^T$

Metric of Performance

- Binary classification: Confusion matrix

		Predicted condition	
		Positive (PP)	Negative (PN)
Actual condition	Total population = P + N		
	Positive (P)	True positive (TP)	False negative (FN)
	Negative (N)	False positive (FP)	True negative (TN)

- Recall: $\frac{TP}{P} = \frac{TP}{TP+FN}$
亦稱：sensitivity; hit rate, true positive rate
- Precision $\frac{TP}{TP+FP}$
- Accuracy = $\frac{TP+TN}{TP+FN+TN+FP}$

F scores

$$F_{\beta} = (1 + \beta^2) \frac{p \cdot r}{\beta^2 p + r}$$

Proof.

$$\frac{\partial F_{\beta}}{\partial \beta} = \frac{2\beta pr^2}{(\beta^2 p + r)^2} > 0 \quad \text{if } \beta > 0.$$

Therefore F_{β} is increasing. Furthermore, if $r > p$, then

$$\frac{\beta^2}{r} + \frac{1}{r} < \frac{\beta^2}{r} + \frac{1}{p} < \frac{\beta^2}{p} + \frac{1}{p}$$

Thus their inverse have the following relation

$$r = \frac{1 + \beta^2}{\frac{\beta^2}{r} + \frac{1}{r}} > \frac{1 + \beta^2}{\frac{\beta^2}{r} + \frac{1}{p}} = F_{\beta} > \frac{1 + \beta^2}{\frac{\beta^2}{p} + \frac{1}{p}} = p.$$

也就是 $p < F_{1/2} < F_1 < F_2 < \dots < r$



Multiple Classification

Confusion matrix

	P ₁	P ₂	P ₃
T ₁	a	b	c
T ₂	d	e	f
T ₃	g	h	i

- 若有 none，假設 class 1=none

- ▶ Recall $r = \frac{e+f}{d+e+f+g+h+i}$

- ▶ Precision $p = \frac{e+f}{b+e+h+c+f+i}$

- 若沒有 none，分別統計

$$r_1 = \frac{a}{a+b+c}, p_1 = \frac{a}{a+d+g}, r_2 = \frac{e}{d+e+f}, p_2 = \frac{e}{b+e+h}, r_3 = \frac{i}{g+h+i},$$

$$p_3 = \frac{i}{c+f+i}$$

F scores for multiple classifications

Compute

$$p_i = \frac{TP_i}{TP_i + FP_i} \quad r_i = \frac{TP_i}{TP_i + FN_i}$$

for each class

- Micro F1:

$$\text{precision } p = \frac{\sum_{i=1}^C TP_i}{\sum_{i=1}^C TP_i + \sum_{i=1}^C FP_i}, \text{ and recall } r = \frac{\sum_{i=1}^C TP_i}{\sum_{i=1}^C TP_i + \sum_{i=1}^C FN_i}.$$

Notice that $p = r$. Micro $F_1 = p = r$

- F scores for each class

$$F_{\beta,i} = (1 + \beta^2) \frac{p_i \cdot r_i}{\beta^2 p_i + r_i}$$

$$\text{Macro } F_{\beta} = \frac{1}{C} \sum_{i=1}^C F_{\beta,i}$$

- 類別不平衡時，Macro F scores 較公平

- Accuracy = $\frac{\sum_{i=1}^C TP_i}{N}$

Good metric?

評估指標依產業而不同

- $F_1 = \frac{2}{\frac{1}{r} + \frac{1}{p}}$ harmonic mean, 均衡 recall 與 precision. 一般分類任務採用 F_1
- Image classification: The top-1, top-k error
Top-1 error = 沒猜中的比例 (1-accuracy), Top-5 error = GT 不在前五名猜測範圍.
- 與安全性相關的產業: 醫療、指紋辨識、人臉辨識解鎖
高 precision
- 廣告、關鍵字
高 recall
- 翻譯任務:
 - ▶ Perplexity $2^{-\ell}$, $\ell = \frac{1}{N} \sum_{i=1}^N \log(w_i | w_{i-1})$

機器翻譯指標: BLEU (Bilingual Evaluation Understudy)

score



$$BLEU = BP \cdot (\prod_{i=1}^4 p_i)^{1/4},$$

where

$$p_i = \frac{\sum_{snt \in \text{Cand-Corpus}} \sum_{i \in snt} \min(m_{cand}^i, m_{ref}^i)}{w_t = \sum_{snt' \in \text{Cand-Corpus}} \sum_{i' \in snt'} \min(m_{snt'}^i, m_{cand}^i)}$$

- m_{cand}^i is the count of i-gram in candidate matching the reference translation
- m_{ref}^i is the count of i-gram in the reference translation
- w_t is the total number of i-grams in candidate translation
- $BP = \min(1, \exp(1 - rl/ol))$: Brevity penalty (懲罰因子)
ol(output-length = 機器譯文長度, rl(reference-length) = 參考翻譯長度)
 - ▶ BLEU 需要計算翻譯後 p_n : n-gram 精確率
 - ▶ BP: 若翻譯後長度小於參考譯文, 則 BP 小於 1
 - ▶ BLEU 的 1-gram 精確率表示譯文終於原文 (loss), 其他 n-gram 表示翻譯流暢程度

假設

Reference: the cat is on the mat

Candidate: the the the cat mat (機器翻譯)

Unigram	m_{cand}^1	m_{ref}^1	$\min(m_{cand}^1, m_{ref}^1)$
the	3	2	2
cat	1	1	1
is	0	1	0
on	0	1	0
mat	1	1	1

$$w_t^1 = 5, p_1 = 2 + 1 + 1/5 = 0.8, BP = \min(1, e^{1-6/5}) = e^{-0.2}$$

2-gram precision

2-gram	m_{cand}^2	m_{ref}^2	$\min(m_{cand}^2, m_{ref}^2)$
the cat	1	1	1
cat is	0	1	0
is on	0	1	0
on the	0	1	0
the mat	0	1	0

$$w_t^2 = 5, p_2 = 1/5 = 0.2$$

3-gram precision

3-gram	m_{cand}^3	m_{ref}^3	$\min(m_{cand}^3, m_{ref}^3)$
the cat is	0	1	0
cat is on	0	1	0
is on the	0	1	0
on the mat	0	1	0

$$w_t^3 = 4, p_3 = 0/5 = 0$$

$$\text{BLEU} = 0$$

作業五

Calculating the BLEU score

Reference: The NASA Opportunity rover is battling a massive dust storm on Mars.

Candidate 1: The Opportunity rover is combating a big sandstorm on Mars.

Candidate 2: A NASA rover is fighting a massive storm on Mars.

作業六

- 找一個 dataset 打造一個以 RNN 為基礎，做情意分析
- 炎龍老師 Mooc 課程 <https://github.com/yenlung/Deep-Learning-MOOC/blob/master/04-1.%20RNN.ipynb>
- 選一個自己喜歡的 dataset，或者IMDB
- 觀察 dataset，正樣本多？還是負樣本多？有無資料不平衡問題？如何解決？
- 模型：RNN, LSTM, BERT, GPT? 幾層？hidden dimension 多少才適合？
- Loss 的選擇
- 訓練到什麼時候才算訓練好？評判指標是什麼
- 準確率如何？採用何種指標？是適合這個 dataset 的評判嗎？

Binary Cross-Entropy

The binary cross entropy is

$$L = -(y \log p + (1 - y) \log(1 - p))$$

where $p \in [0, 1] \subset \mathbb{R}^1$ and $y \in \{0, 1\}$.

Categorical cross-entropy is given by



$$L = -\mathbf{y} \log \mathbf{p}$$

Suppose there are two categories. That is, $\mathbf{y}, \mathbf{p} \in \mathbb{R}^2$, where $\mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \end{pmatrix}$

and $\mathbf{p} = \begin{pmatrix} p_1 \\ p_2 \end{pmatrix}$. Since \mathbf{y} is the ground truth and \mathbf{p} represents the probability of predicts, we have $y_1 + y_2 = 1$ and $p_1 + p_2 = 1$. Then

$$L = -\mathbf{y} \log \mathbf{p} = -(y_1 \log p_1 + y_2 \log p_2) = -(y_1 \log p_1 + (1 - y_1) \log(1 - p_1))$$


Loss



Focal Loss


論文：Focal Loss for Dense Object Detection

$$FL = -\frac{1}{n} \sum_{i=1}^n y_i (1 - p_i)^\gamma \log(p_i)$$



	分錯	分對但預測差	分很對
Cross Entropy	$\frac{L_{CE}(0.4)}{L_{CE}(0.8)} = 4.11$	$\frac{L_{CE}(0.5)}{L_{CE}(0.9)} = 6.58$	$\frac{L_{CE}(0.8)}{L_{CE}(0.9)} = 2.12$
Focal Loss(1)	$\frac{L_{FL1}(0.4)}{L_{FL1}(0.8)} = 12.32$	$\frac{L_{FL1}(0.5)}{L_{FL1}(0.9)} = 32.89$	$\frac{L_{FL1}(0.8)}{L_{FL1}(0.9)} = 4.24$
Focal Loss(2)	$\frac{L_{FL2}(0.4)}{L_{FL2}(0.8)} = 36.96$	$\frac{L_{FL2}(0.5)}{L_{FL2}(0.9)} = 164.47$	$\frac{L_{FL2}(0.8)}{L_{FL2}(0.9)} = 8.47$

實際值 $y = 0$ 時：



	分錯	分對但預測差	分很對
Cross Entropy	$\frac{L_{CE}(0.6)}{L_{CE}(0.2)} = 4.11$	$\frac{L_{CE}(0.49)}{L_{CE}(0.1)} = 6.39$	$\frac{L_{CE}(0.2)}{L_{CE}(0.1)} = 2.12$
Focal Loss(1)	$\frac{L_{FL1}(0.6)}{L_{FL1}(0.2)} = 12.32$	$\frac{L_{FL1}(0.49)}{L_{FL1}(0.1)} = 31.32$	$\frac{L_{FL1}(0.2)}{L_{FL1}(0.1)} = 4.24$
Focal Loss(2)	$\frac{L_{FL2}(0.6)}{L_{FL2}(0.2)} = 36.96$	$\frac{L_{FL2}(0.5)}{L_{FL2}(0.9)} = 153.44$	$\frac{L_{FL2}(0.2)}{L_{FL2}(0.1)} = 8.47$

Focal Loss 特色

- 著重在處理預測差的問題 (hard, misclassified examples)
- 類別不均衡問題：通常 object detection 採用 OHEH
- α -balanced variant

$$FL(p_t) = -\alpha(1 - p_t)^\gamma \log(p_t)$$

default $\gamma = 2$, $\alpha = 0.25$ (γ 增加, α 減少). In practice, α may be set by inverse class frequency or treated as a hyperparameter to set by cross-validation.